

AI Operations Budget - Cap Revision Request

Prepared by: Ken Mun (CTO) | Submitted to: Angie Foong (CEO) | 28 April 2026

THE ASK: A\$500/month -> A\$1,500/month for 3 months

The original cap was set before live data existed. 4 days of operation shows spend is 4-5x the assumption. This reflects an infrastructure build phase - not waste. Optimisations already in place project steady-state at A\$900-1,100/month. A\$1,500 provides a safe ceiling through OC2 deployment.

SECTION 1

What the Data Shows - 4 Days Live

\$169 USD 4-DAY TOTAL SPEND	A\$2,100 MONTH RUN-RATE (RAW)	A\$500 CURRENT CAP/MONTH	4.2x GAP VS ASSUMPTION
--	--	---	---

Day	Date	USD Spend	Turns	Phase	Note
Day 1	25 Apr	\$25.75	416	Build	Foundation: agents, identity, scripts, gateway
Day 2	26 Apr	\$37.15	499	Build	Resiliency stack, ITSM, PVT, diagnostics
Day 3	27 Apr	\$90.07	793	OUTLIER	Aria + governance launch + Opus model drift ~4hrs at 3x cost. Non-recurring.
Day 4	28 Apr	~\$22 proj.	200+	Optimised	3-tier model strategy live. Haiku on sub-tasks.

CONTEXT ON DAY 3 - WHY IT WAS AN OUTLIER

Day 3 combined: launching Aria (second agent), building the entire governance layer (4 agents), and two model drift incidents where the wrong model ran for ~4 hours at 3x cost. All three drivers are fully resolved and non-recurring. Warden now audits all 6 agents every 15 minutes.

SECTION 2

What A\$260 Bought in 4 Days

For context: a senior developer day in Melbourne costs A\$800-1,200. 4 AI-accelerated days at ~A\$65/day delivered:

Capability	Status	Business Value
Yoda - Lead Ops Agent	Live	Runs operations 24/7. Stand-up, auto-heal, diagnostics, git, Notion sync, incident response
Aria - Business Lead Agent	Live	Angie's AI partner for business stream, onboarding pipeline and training content.
Governance (Shield/Lex/Sage)	Layer Live	Every external output reviewed for security, legal and quality before it leaves.
Warden - Model Compliance	Live	Audits all 6 agents every 15 min. Catches model drift before it costs money.
Resiliency Stack	Live	Auto-heal, diagnostics, audit trail, anti-drift. Platform self-monitors and self-repairs overnight
3-Tier Model Strategy	Live	Routing engine: right model for each task. Already reducing sub-task costs by 3x.

DEVELOPER COST EQUIVALENT

15-20 days of senior developer work delivered in 4 days at A\$260 in AI API cost. Equivalent developer cost at market rate: A\$12,000-24,000. Before counting 24/7 autonomous operations and ongoing monitoring.

AI Operations Budget - Cap Revision Request

Prepared by: Ken Mun (CTO) | Submitted to: Angie Foong (CEO) | 28 April 2026

SECTION 3

Phase	Timeframe	USD/day	AUD/month	Key Driver
Build phase	Now - Jul 2026 (OC2)	\$28-35	A\$1,320-1,650	Active dev, new integrations, agents
Steady state	Post-OC2	\$15-22	A\$700-1,030	Ops only + Gemma4:26b delegation
Original estimate	Set Apr 2026	\$4	A\$500	reducing API reliance. No agent infrastructure costed.

WHY THE ORIGINAL ESTIMATE WAS WRONG

The A\$500 cap was set before Day 1 - before knowing how many agents, crons, governance reviews, and interactive sessions the platform would run. It was a placeholder, not a costed estimate. With 4 days of real data and a live cost tracker, we can now plan accurately.

SECTION 4

Assumptions

- > AUD/USD rate held at 0.65 for the duration.
- > Claude Sonnet 4.6 remains primary interactive model (\$3/\$15 per MTok).
- > Haiku 4.5 routing active - reduces sub-task cost by ~3x on ~30 turns/day.
- > gemma4:e2b handles all background crons at zero API cost. Already live.
- > OC2 (32GB Mac mini) deployed by July 2026 - enables Gemma4:26b delegation, reducing API reliance.
- > No major new agent builds after June. Incremental improvements only.
- > API pricing does not increase significantly.
- > No Opus usage unless explicitly triggered by Ken for high-stakes tasks.

SECTION 5

Controls Already in Place

Control	How it works	Status	Effective
Daily burn alert	Telegram alert to Ken if daily spend >= \$40 USD. Fires	Live	Today
3-tier routing	Right model for each task automatically. Haiku for	Live	Today
Warden drift checks	sub-tasks. Gemma4 for crons. Every 15 min catches if a cheap model drifts to	Live	Day 4
3-tier credit alerts	expensive. Escalates to Ken. API balance alerts at \$50, \$25, \$10. Both Ken and Angie	Live	Day 3
Monthly model review	notified. 28th each month - full cost audit + benchmark. Ken	Live	Day 4
Delegation logging	sign-off required to change anything. Every routed sub-task logged. Model, latency, outcome, savings estimate.	Live	Today

RECOMMENDATION

A\$1,500	3 months	A\$4,500	Jul 2026
NEW MONTHLY CAP	DURATION (BUILD PHASE)	MAX EXPOSURE	REVIEW AT OC2

A\$1,500/month sits at the midpoint of the build-phase projection (A\$1,320-1,650). It provides a realistic ceiling without slowing the build. At OC2 deployment the cap should be reviewed downward to ~A\$1,000/month as Gemma4:26b delegation reduces API reliance. The A\$4,500 3-month total compares against A\$30,000-50,000 to hire a developer for equivalent infrastructure - a 7-11x cost advantage, before counting 24/7 autonomous operations.

AI Operations Budget - Cap Revision Request

Prepared by: Ken Mun (CTO) | Submitted to: Angie Foong (CEO) | 28 April 2026

DECISION & SIGN-OFF

APPROVED BY (CEO)

SUBMITTED BY (CTO)

DATE

Angie Foong - CEO, Alanchors

Ken Mun - CTO, Alanchors